

## Question 5

**Max accuracy** was achieved when I used the following hyperparameters:

- `n_estimators=700`
- `min_samples_leaf=2`
- `max_depth=4`
- `max_features=None`
- `learning_rate=0.1`

For **best** implementation:

```
Accuracy: 0.9974782852339591  
Precision: 0.997039230199852  
Recall: 1.0
```

*As `n_estimators` increases, all three (accuracy, precision, recall) seem to improve. Gradient boosting is robust to overfitting, and hence, it performs better as we increase `n_estimators`.*

For `n_estimators=100`:

```
Accuracy: 0.9948164752031381  
Precision: 0.9940949725252194  
Recall: 0.9998350243339107
```

For `n_estimators=300`:

```
Accuracy: 0.9965676660128888  
Precision: 0.996137409598948  
Recall: 0.9998350243339107
```

For `n_estimators=500`:

```
Accuracy: 0.997128047072009  
Precision: 0.9967927631578948  
Recall: 0.9998350243339107
```

For n\_estimators=700:

Accuracy: 0.997338189969179

Precision: 0.9969569865942923

Recall: 0.9999175121669553

For a **simple decision tree** implementation:

Accuracy: 0.9912440459512468

Precision: 0.9952121512299819

Recall: 0.99447331518601

*Clearly, gradient boosting is much better than a simple decision tree in terms of accuracy. Precision, recall too, are better for boosting method.*